

Ex no: 5

Create tables in Hive and write queries to access the data in the table

Aim:

To create tables in Hive and write queries to access the data in the table.

Procedure:

1. Install Apache Derby version 10.14.2.0.
2. Install Apache Hive version 3.1.2.
3. Set up environment variables for Hive and Derby.
4. Copy all the jar files in the “db-derby-10.14.2.0-bin\lib” folder and paste it in the Hive libraries directory “apache-hive-3.1.3-bin\lib”.
5. Start Hadoop services
6. Run the command “StartNetworkServer -h 0.0.0.0” in a new command window to start Derby.
7. To start Hive run the following command “hive --service schematool -dbType derby -initSchema”
8. After this open the Hive shell by typing the command “hive”.
9. To create a new database type the command “CREATE DATABASE sample;”.
10. To verify if the database has been create type the command “SHOW DATABASES;”.
11. To use the database type “USE sample;”
12. Now, create a table in the database using the command “CREATE TABLE students (name STRING, roll INT, dept STRING);”
13. To add data to the table from a csv file in your local device to the hive database use the command “LOAD DATA LOCAL INPATH ‘C:\EX_5_Hive\students.csv’ INTO students;”.
14. Type the command “SHOW TABLES;” to display the tables in the currently using database.
15. To display the table’s description type the command “DESC students;”.
16. To display all the data in the table, type the command “SELECT * FROM students;”

Output Screenshots:

DERBY_HOME	C:\db-derby-10.14.2.0-bin
HADOOP_HOME	C:\hadoop-3.3.6
HIVE_BIN	%HIVE_HOME%\bin
HIVE_HOME	C:\apache-hive-3.1.3-bin
HIVE_LIB	%HIVE_HOME%\lib
JAVA_HOME	C:\Program Files\Java\jdk1.8.0_202\bin

%HIVE_BIN%

%DERBY_HOME%\bin

lib

This PC > Windows (C:) > db-derby-10.14.2.0-bin > lib

Search lib

New Sort View Details

Name	Date modified	Type	Size
derby	07-04-2018 06:40	Executable Jar File	3,158 KB
derby.war	07-04-2018 06:40	WAR File	2 KB
derbyclient	07-04-2018 06:40	Executable Jar File	575 KB
derbyLocale_cs	07-04-2018 06:40	Executable Jar File	93 KB
derbyLocale_de_DE	07-04-2018 06:40	Executable Jar File	110 KB
derbyLocale_es	07-04-2018 06:40	Executable Jar File	104 KB
derbyLocale_fr	07-04-2018 06:40	Executable Jar File	110 KB
derbyLocale_hu	07-04-2018 06:40	Executable Jar File	94 KB
derbyLocale_it	07-04-2018 06:40	Executable Jar File	104 KB
derbyLocale_ja_JP	07-04-2018 06:40	Executable Jar File	121 KB
derbyLocale_ko_KR	07-04-2018 06:40	Executable Jar File	115 KB
derbyLocale_pl	07-04-2018 06:40	Executable Jar File	92 KB
derbyLocale_pt_BR	07-04-2018 06:40	Executable Jar File	89 KB
derbyLocale_ru	07-04-2018 06:40	Executable Jar File	119 KB
derbyLocale_zh_CN	07-04-2018 06:40	Executable Jar File	108 KB
derbyLocale_zh_TW	07-04-2018 06:40	Executable Jar File	109 KB
derbynet	07-04-2018 06:40	Executable Jar File	267 KB

20 items

lib

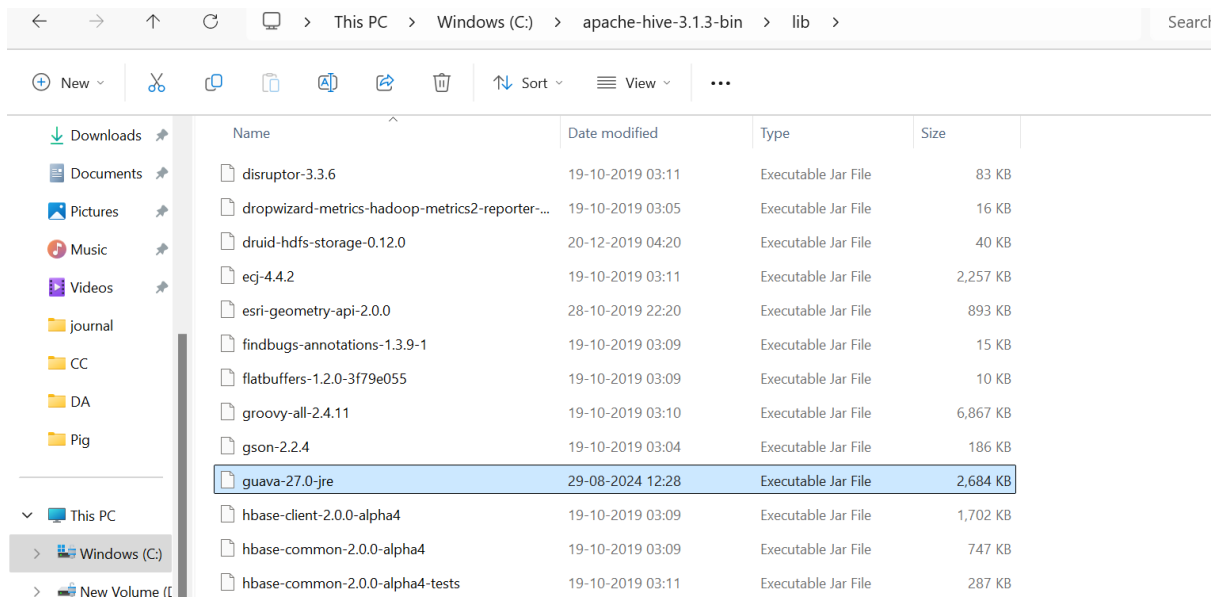
This PC > Windows (C:) > apache-hive-3.13-bin > lib

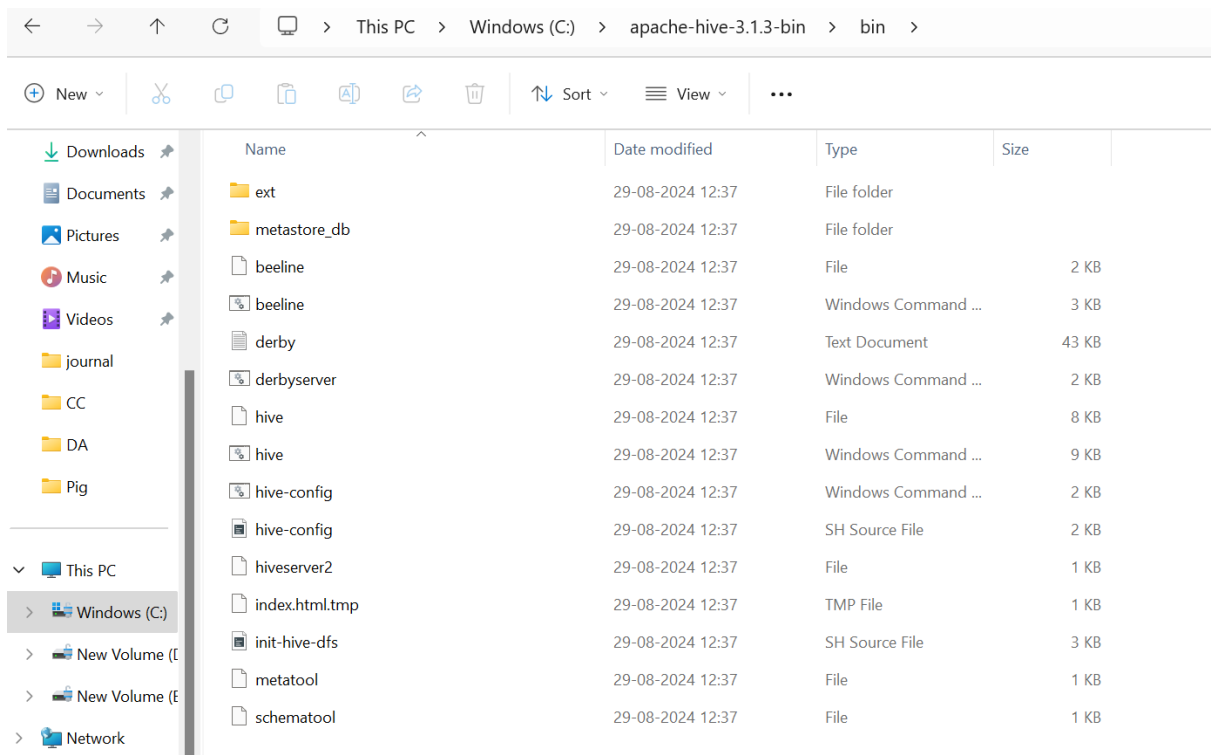
Search lib

New Sort View Details

Name	Date modified	Type	Size
commons-vfs2-2.1	19-10-2019 03:11	Executable Jar File	432 KB
compress-lzf-1.0.3	19-10-2019 03:09	Executable Jar File	78 KB
curator-client-2.12.0	28-10-2019 21:06	Executable Jar File	2,367 KB
curator-framework-2.12.0	28-10-2019 21:05	Executable Jar File	198 KB
curator-recipes-2.12.0	28-10-2019 21:05	Executable Jar File	277 KB
datanucleus-api-jdo-4.2.4	19-10-2019 03:10	Executable Jar File	359 KB
datanucleus-core-4.1.17	19-10-2019 03:10	Executable Jar File	1,970 KB
datanucleus-rdbms-4.1.19	19-10-2019 03:10	Executable Jar File	1,864 KB
derby	07-04-2018 06:40	Executable Jar File	3,158 KB
derby.war	07-04-2018 06:40	WAR File	2 KB
derby-10.14.1.0	19-10-2019 03:05	Executable Jar File	3,157 KB
derbyclient	07-04-2018 06:40	Executable Jar File	575 KB
derbyLocale_cs	07-04-2018 06:40	Executable Jar File	93 KB
derbyLocale_de_DE	07-04-2018 06:40	Executable Jar File	110 KB
derbyLocale_es	07-04-2018 06:40	Executable Jar File	104 KB
derbyLocale_fr	07-04-2018 06:40	Executable Jar File	110 KB
derbyLocale_hu	07-04-2018 06:40	Executable Jar File	94 KB

279 items | 26 items selected 11.8 MB





```
Administrator: Command Prompt - hive
Microsoft Windows [Version 10.0.22631.4112]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd /

C:\>start-all
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\>jps
9600 NameNode
14068 DataNode
18100 NodeManager
16184 ResourceManager
2600 Jps
```

```
Administrator: Command Prompt - StartNetworkServer -h 0.0.0.0
Microsoft Windows [Version 10.0.22631.4112]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd /

C:\>StartNetworkServer -h 0.0.0.0
Thu Aug 29 12:58:55 IST 2024 : Security manager installed using the Basic server security policy.
Thu Aug 29 12:58:55 IST 2024 : Apache Derby Network Server - 10.14.2.0 - (1828579) started and ready to accept connections on port 1527
```

```
Administrator: Command Prompt - hive
C:\>hive --service schematool -dbType derby -initSchema
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/C:/hadoop-3.3.6/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/C:/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Reload4jLoggerFactory]
2024-08-29 13:07:05,856 INFO conf.HiveConf: Found configuration file null
2024-08-29 13:07:06,164 INFO tools.HiveSchemaHelper: Metastore connection URL: jdbc:derby;;databaseName=metastore_db;create=true
Metastore connection URL: jdbc:derby;;databaseName=metastore_db;create=true
2024-08-29 13:07:06,166 INFO tools.HiveSchemaHelper: Metastore Connection Driver : org.apache.derby.jdbc.EmbeddedDriver
Metastore Connection Driver : org.apache.derby.jdbc.EmbeddedDriver
2024-08-29 13:07:06,166 INFO tools.HiveSchemaHelper: Metastore connection User: APP
Metastore connection User: APP
Starting metastore schema initialization to 3.1.0
Initialization script hive-schema-3.1.0.derby.sql
```

```
Administrator: Command Prompt - hive

Initialization script completed
schemaTool completed

C:\>hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/C:/hadoop-3.3.6/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/C:/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Reload4jLoggerFactory]
2024-08-29 13:07:29,882 INFO conf.HiveConf: Found configuration file null
2024-08-29 13:07:30,957 WARN common.LogUtils: hive-site.xml not found on CLASSPATH
Hive Session ID = ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:07:31,003 INFO SessionState: Hive Session ID = ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9

Logging initialized using configuration in jar:file:/C:/apache-hive-3.1.3-bin/lib/hive-common-3.1.3.jar!/hive-log4j2.properties Async: true
2024-08-29 13:07:31,053 INFO SessionState:
Logging initialized using configuration in jar:file:/C:/apache-hive-3.1.3-bin/lib/hive-common-3.1.3.jar!/hive-log4j2.properties Async: true
2024-08-29 13:07:31,911 INFO session.SessionState: Created HDFS directory: /tmp/hive/mannu
```

```
2024-08-29 13:07:40,952 INFO metastore.HiveMetaStore: 1: get_multi_table : db=default tbls=
2024-08-29 13:07:40,953 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_multi_table : db=default tbls=
2024-08-29 13:07:40,956 INFO metadata.HiveMaterializedViewsRegistry: Materialized views registry has been initialized
hive>
```

```
Administrator: Command Prompt - hive

hive> CREATE DATABASE sample;
2024-08-29 13:17:58,469 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:17:58,603 INFO ql.Driver: Compiling command(queryId=mannu_20240829131758_f94713b1-6ff5-4cb0-b31c-56a2aa75ce78): CREATE DATABASE sample
2024-08-29 13:17:58,907 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:17:58,923 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-08-29 13:17:58,926 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-08-29 13:17:58,932 INFO ql.Driver: Completed compiling command(queryId=mannu_20240829131758_f94713b1-6ff5-4cb0-b31c-56a2aa75ce78); Time taken: 0.375 seconds
2024-08-29 13:17:58,932 INFO reexec.ReExecDriver: Execution #1 of query
2024-08-29 13:17:58,933 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:17:58,934 INFO ql.Driver: Executing command(queryId=mannu_20240829131758_f94713b1-6ff5-4cb0-b31c-56a2aa75ce78): CREATE DATABASE sample
2024-08-29 13:17:58,945 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-29 13:17:58,965 INFO sqlstd.SQLStdHiveAccessController: Created SQLStdHiveAccessController for session context : HiveAuthzSessionContext [sessionString=ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9, clientType=HIVECLI]
2024-08-29 13:17:58,968 WARN session.SessionState: METASTORE_FILTER_HOOK will be ignored, since hive.security.authorization.manager is set to instance of HiveAuthorizerFactory.
2024-08-29 13:17:58,969 INFO metastore.HiveMetaStoreClient: Metastore configuration metastore.filter.hook changed from org.apache.hadoop.hive.metastore.DefaultMetaStoreFilterHookImpl to org.apache.hadoop.hive.ql.security.authorization.plugin.AuthorizationMetaStoreFilterHook
2024-08-29 13:17:58,971 INFO metastore.HiveMetaStore: 0: Cleaning up thread local RawStore...
```

```
Administrator: Command Prompt - hive
126): SHOW DATABASES
2024-08-29 13:20:43,574 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-29 13:20:43,575 INFO metastore.HiveMetaStore: 0: get_databases: @hive#
2024-08-29 13:20:43,575 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_databases: @hive#

2024-08-29 13:20:43,583 INFO exec.DDLTask: results : 2
2024-08-29 13:20:43,619 INFO ql.Driver: Completed executing command(queryId=mannu_20240829132043_8f193ffe-60bf-411c-baa8-84d34fe34126); Time taken: 0.045 seconds
OK
2024-08-29 13:20:43,620 INFO ql.Driver: OK
2024-08-29 13:20:43,620 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:20:43,631 INFO Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
2024-08-29 13:20:43,669 INFO mapred.FileInputFormat: Total input files to process : 1
2024-08-29 13:20:43,697 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:2,
default
sample
Time taken: 0.234 seconds, Fetched: 2 row(s)
2024-08-29 13:20:43,707 INFO CliDriver: Time taken: 0.234 seconds, Fetched: 2 row(s)
2024-08-29 13:20:43,707 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:20:43,708 INFO session.SessionState: Resetting thread name to main
hive>
```

```
Administrator: Command Prompt - hive
b44): USE sample
2024-08-29 13:26:07,735 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-29 13:26:07,736 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-29 13:26:07,736 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_database: @hive#sample

2024-08-29 13:26:07,738 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-29 13:26:07,738 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_database: @hive#sample

2024-08-29 13:26:07,740 INFO ql.Driver: Completed executing command(queryId=mannu_20240829132607_94ec5c2e-2f90-4aea-8260-79d23eb2eb44); Time taken: 0.007 seconds
OK
2024-08-29 13:26:07,741 INFO ql.Driver: OK
2024-08-29 13:26:07,742 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
Time taken: 0.068 seconds
2024-08-29 13:26:07,743 INFO CliDriver: Time taken: 0.068 seconds
2024-08-29 13:26:07,747 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:26:07,748 INFO session.SessionState: Resetting thread name to main
hive>
```

```
Administrator: Command Prompt - hive
hive> CREATE TABLE students (name STRING, roll INT, dept STRING);
2024-08-29 13:28:32,064 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:28:32,065 INFO session.SessionState: Updating thread name to ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9 main
2024-08-29 13:28:32,067 INFO ql.Driver: Compiling command(queryId=mannu_20240829132832_025f466a-0d8e-4e9f-9e39-0ba100314c7b): CREATE TABLE students (name STRING, roll INT, dept STRING)
2024-08-29 13:28:32,128 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:28:32,135 INFO parse.CalcitePlanner: Starting Semantic Analysis
2024-08-29 13:28:32,149 INFO parse.CalcitePlanner: Creating table sample.students position=13
2024-08-29 13:28:32,163 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-29 13:28:32,164 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_database: @hive#sample

2024-08-29 13:28:32,181 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-08-29 13:28:32,181 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-08-29 13:28:32,182 INFO ql.Driver: Completed compiling command(queryId=mannu_20240829132832_025f466a-0d8e-4e9f-9e39-0ba100314c7b); Time taken: 0.115 seconds
2024-08-29 13:28:32,182 INFO reexec.ReExecDriver: Execution #1 of query
2024-08-29 13:28:32,183 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:28:32,183 INFO ql.Driver: Executing command(queryId=mannu_20240829132832_025f466a-0d8e-4e9f-9e39-0ba100314c7b): CREATE TABLE students (name STRING, roll INT, dept STRING)
2024-08-29 13:28:32,196 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-29 13:28:32,426 INFO metastore.HiveMetaStore: 0: create table: Table(tableName=students, dbName=sample, owner=mannu, createTime:1724918312, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name=name, type:string
```



```
Administrator: Command Prompt - hive
hive> LOAD DATA LOCAL INPATH 'C:\EX_5_Hive\students.csv' INTO TABLE students;
2024-08-29 14:11:39,111 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 14:11:39,111 INFO session.SessionState: Updating thread name to ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9 main
2024-08-29 14:11:39,114 INFO ql.Driver: Compiling command(queryId=mannu_20240829141139_a912a1e1-e840-4b7e-93d2-09ee931597aa): LOAD DATA LOCAL INPATH 'C:\EX_5_Hive\students.csv' INTO TABLE students
2024-08-29 14:11:39,138 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 14:11:39,138 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-29 14:11:39,140 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-29 14:11:39,161 INFO io.HiveFileFormatUtils: Failed to check mode for C:\EX_5_Hive\students.csv: View 'unix' not available (class java.lang.UnsupportedOperationException)
2024-08-29 14:11:39,198 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-08-29 14:11:39,199 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-08-29 14:11:39,200 INFO ql.Driver: Completed compiling command(queryId=mannu_20240829141139_a912a1e1-e840-4b7e-93d2-09ee931597aa); Time taken: 0.086 seconds
2024-08-29 14:11:39,200 INFO reexec.ReExecDriver: Execution #1 of query
2024-08-29 14:11:39,201 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 14:11:39,201 INFO ql.Driver: Executing command(queryId=mannu_20240829141139_a912a1e1-e840-4b7e-93d2-09ee931597aa): LOAD DATA LOCAL INPATH 'C:\EX_5_Hive\students.csv' INTO TABLE students
2024-08-29 14:11:39,202 INFO ql.Driver: Starting task [Stage-0:MOVE] in serial mode
2024-08-29 14:11:39,203 INFO metastore.HiveMetaStore: 0: Cleaning up thread local RawStore...
2024-08-29 14:11:39,203 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=Cleaning up thread local RawStore...
2024-08-29 14:11:39,204 INFO metastore.HiveMetaStore: 0: Done cleaning up thread local RawStore
2024-08-29 14:11:39,204 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=Done cleaning up thread local RawStore
Loading data to table sample.students
2024-08-29 14:11:39,205 INFO exec.Task: Loading data to table sample.students from file:C:\EX_5_Hive\students.csv
2024-08-29 14:11:39,208 INFO metastore.HiveMetaStore: 0: Opening raw store with implementation class:org.apache.hadoop.hive.metastore.ObjectStore
2024-08-29 14:11:39,209 WARN metastore.ObjectStore: datanucleus.autoStartMechanismMode is set to unsupported value null . Setting it to value: ignored
2024-08-29 14:11:39,209 INFO metastore.ObjectStore: ObjectStore, initialize called
2024-08-29 14:11:39,212 INFO metastore.MetaStoreDirectSql: Using direct SQL, underlying DB is DERBY
```

```
OK
2024-08-29 14:11:40,122 INFO ql.Driver: OK
2024-08-29 14:11:40,122 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
Time taken: 1.008 seconds
2024-08-29 14:11:40,123 INFO CliDriver: Time taken: 1.008 seconds
2024-08-29 14:11:40,123 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 14:11:40,124 INFO session.SessionState: Resetting thread name to main
```

```
Administrator: Command Prompt - hive
873): SHOW TABLES
2024-08-29 13:38:31,138 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-29 13:38:31,139 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-29 13:38:31,142 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_database: @hive#sample
2024-08-29 13:38:31,145 INFO metastore.HiveMetaStore: 0: get_tables: db=@hive#sample pat=.*
2024-08-29 13:38:31,146 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_tables: db=@hive#sample pat=.*
2024-08-29 13:38:31,165 INFO ql.Driver: Completed executing command(queryId=mannu_20240829133831_b2cef26b-7f7e-480b-af4f-5415a2e61873); Time taken: 0.029 seconds
OK
2024-08-29 13:38:31,166 INFO ql.Driver: OK
2024-08-29 13:38:31,166 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:38:31,170 INFO mapred.FileInputFormat: Total input files to process : 1
2024-08-29 13:38:31,180 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:1, students
Time taken: 0.069 seconds, Fetched: 1 row(s)
2024-08-29 13:38:31,184 INFO CliDriver: Time taken: 0.069 seconds, Fetched: 1 row(s)
2024-08-29 13:38:31,184 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:38:31,185 INFO session.SessionState: Resetting thread name to main
hive>
```

```
Administrator: Command Prompt - hive
858): DESC students
2024-08-29 13:40:43,654 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-29 13:40:43,654 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-29 13:40:43,654 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-29 13:40:43,672 INFO ql.Driver: Completed executing command(queryId=mannu_20240829134043_f8d37a98-6cf6-4d4f-b0f6-30bcd856e858); Time taken: 0.019 seconds
OK
2024-08-29 13:40:43,673 INFO ql.Driver: OK
2024-08-29 13:40:43,674 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 13:40:43,677 INFO mapred.FileInputFormat: Total input files to process : 1
2024-08-29 13:40:43,684 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:3, name string
roll int
dept string
Time taken: 0.241 seconds, Fetched: 3 row(s)
2024-08-29 13:40:43,687 INFO CliDriver: Time taken: 0.241 seconds, Fetched: 3 row(s)
2024-08-29 13:40:43,687 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 13:40:43,687 INFO session.SessionState: Resetting thread name to main
hive>
```



```

Administrator: Command Prompt - hive
hive> SELECT * FROM students;
2024-08-29 14:11:47,539 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 14:11:47,539 INFO session.SessionState: Updating thread name to ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9 main
2024-08-29 14:11:47,541 INFO ql.Driver: Compiling command(queryId=mannu_20240829141147_8b78acf5-1036-4f94-9d2f-8c2bd01eecd6): SELECT * FROM students
2024-08-29 14:11:47,562 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 14:11:47,563 INFO parse.CalcitePlanner: Starting Semantic Analysis
2024-08-29 14:11:47,564 INFO parse.CalcitePlanner: Completed phase 1 of Semantic Analysis
2024-08-29 14:11:47,564 INFO parse.CalcitePlanner: Get metadata for source tables
2024-08-29 14:11:47,565 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-29 14:11:47,566 INFO metastore.HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-29 14:11:47,579 INFO parse.CalcitePlanner: Get metadata for subqueries
2024-08-29 14:11:47,579 INFO parse.CalcitePlanner: Get metadata for destination tables
2024-08-29 14:11:47,599 INFO ql.Context: New scratch dir is hdfs://localhost:9000/tmp/hive/mannu/ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9/hive_2024-08-29_14-11-47_561_917565775
703927153-1
2024-08-29 14:11:47,599 INFO parse.CalcitePlanner: Completed getting MetaData in Semantic Analysis
2024-08-29 14:11:47,604 INFO metastore.HiveMetaStore: 0: get_not_null_constraints : tbl=hive.sample.students
2024-08-29 14:11:47,605 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_not_null_constraints : tbl=hive.sample.students
2024-08-29 14:11:47,606 INFO metastore.HiveMetaStore: 0: get_primary_keys : tbl=hive.sample.students
2024-08-29 14:11:47,607 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_primary_keys : tbl=hive.sample.students
2024-08-29 14:11:47,608 INFO metastore.HiveMetaStore: 0: get_primary_keys : tbl=hive.sample.students
2024-08-29 14:11:47,608 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_primary_keys : tbl=hive.sample.students
2024-08-29 14:11:47,610 INFO metastore.HiveMetaStore: 0: get_unique_constraints : tbl=hive.sample.students
2024-08-29 14:11:47,610 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_unique_constraints : tbl=hive.sample.students
2024-08-29 14:11:47,612 INFO metastore.HiveMetaStore: 0: get_foreign_keys : parenttbl=null foreigndb=sample foreigntbl=students
2024-08-29 14:11:47,613 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_foreign_keys : parenttbl=null foreigndb=sample foreigntbl=students
2024-08-29 14:11:47,641 INFO metastore.HiveMetaStore: 0: get_databases: @hive#
2024-08-29 14:11:47,641 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_databases: @hive#
2024-08-29 14:11:47,643 INFO metastore.HiveMetaStore: 0: get_materialized_views_for_rewriting: db=@hive#default
2024-08-29 14:11:47,644 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_materialized_views_for_rewriting: db=@hive#default
2024-08-29 14:11:47,646 INFO metastore.HiveMetaStore: 0: get_materialized_views_for_rewriting: db=@hive#sample
2024-08-29 14:11:47,646 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_materialized_views_for_rewriting: db=@hive#sample
2024-08-29 14:11:47,651 INFO parse.CalcitePlanner: Get metadata for source tables
2024-08-29 14:11:47,651 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-29 14:11:47,651 INFO HiveMetaStore.audit: ugi=mannu ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-29 14:11:47,658 INFO parse.CalcitePlanner: Get metadata for subqueries
2024-08-29 14:11:47,658 INFO parse.CalcitePlanner: Get metadata for destination tables
2024-08-29 14:11:47,661 INFO ql.Context: New scratch dir is hdfs://localhost:9000/tmp/hive/mannu/ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9/hive_2024-08-29_14-11-47_561_917565775

```

```

OK
2024-08-29 14:11:47,682 INFO ql.Driver: OK
2024-08-29 14:11:47,682 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-29 14:11:47,704 INFO mapred.FileInputFormat: Total input files to process : 1
2024-08-29 14:11:47,787 INFO lazy.LazyStruct: Missing fields! Expected 3 fields but only got 1! Ignoring similar problems.
2024-08-29 14:11:47,791 INFO exec.TableScanOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_TS_0:10,
2024-08-29 14:11:47,791 INFO exec.SelectOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_SEL_1:10,
2024-08-29 14:11:47,792 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_3:10,
name,roll,dept NULL NULL
Poppy,110,CSE NULL NULL
Casteel,126,CSE NULL NULL
Sera,231,CSBS NULL NULL
Nyktos,241,CSBS NULL NULL
Sam,343,IT NULL NULL
Emma,356,IT NULL NULL
Jesse,511,ECE NULL NULL
Lily,521,ECE NULL NULL
Kieran,157,CSE NULL NULL
Time taken: 0.141 seconds, Fetched: 10 row(s)
2024-08-29 14:11:47,811 INFO CliDriver: Time taken: 0.141 seconds, Fetched: 10 row(s)
2024-08-29 14:11:47,811 INFO conf.HiveConf: Using the default value passed in for log id: ad7115b2-2eb5-4d3a-8416-b2cbb794b0b9
2024-08-29 14:11:47,812 INFO session.SessionState: Resetting thread name to main
hive>

```

Result:

Thus, the exercise to create tables in Hive and writing queries to access data in the table has been executed successfully.